

# Modern Computer Architecture (Processor Design)

Prof. Dan Connors  
[dconnors@colostate.edu](mailto:dconnors@colostate.edu)

# Computer Architecture

---

- Historic definition

Computer Architecture =

Instruction Set Architecture +

Computer Organization

- Famous architects: Wright, Fuller, Herzog
- Famous computer architects: Smith, Patt, Hwu, Hennessey, Patterson

# Instruction Set Architecture

---

- Important acronym: **ISA**
  - Instruction Set Architecture
- The low-level software interface to the machine
  - Assembly language of the machine
  - Must translate any programming language into this language
  - Examples: IA-32 (Intel instruction set), MIPS, SPARC, Alpha, PA-RISC, PowerPC, ...
- Visible to programmer

# Differences between ISA's

---

- Much more is similar between ISA's than different.  
Compare MIPS & x86:
  - Instructions:
    - same basic types
    - different names and variable-length encodings
    - x86 branches use condition codes (like MIPS floating point)
    - x86 supports (register + memory) -> (register) format
  - Registers:
    - Register-based architecture
    - different number and names, x86 allows partial reads/writes
  - Memory:
    - Byte addressable, 32-bit address space
    - x86 has additional addressing modes and many instruction types can access memory

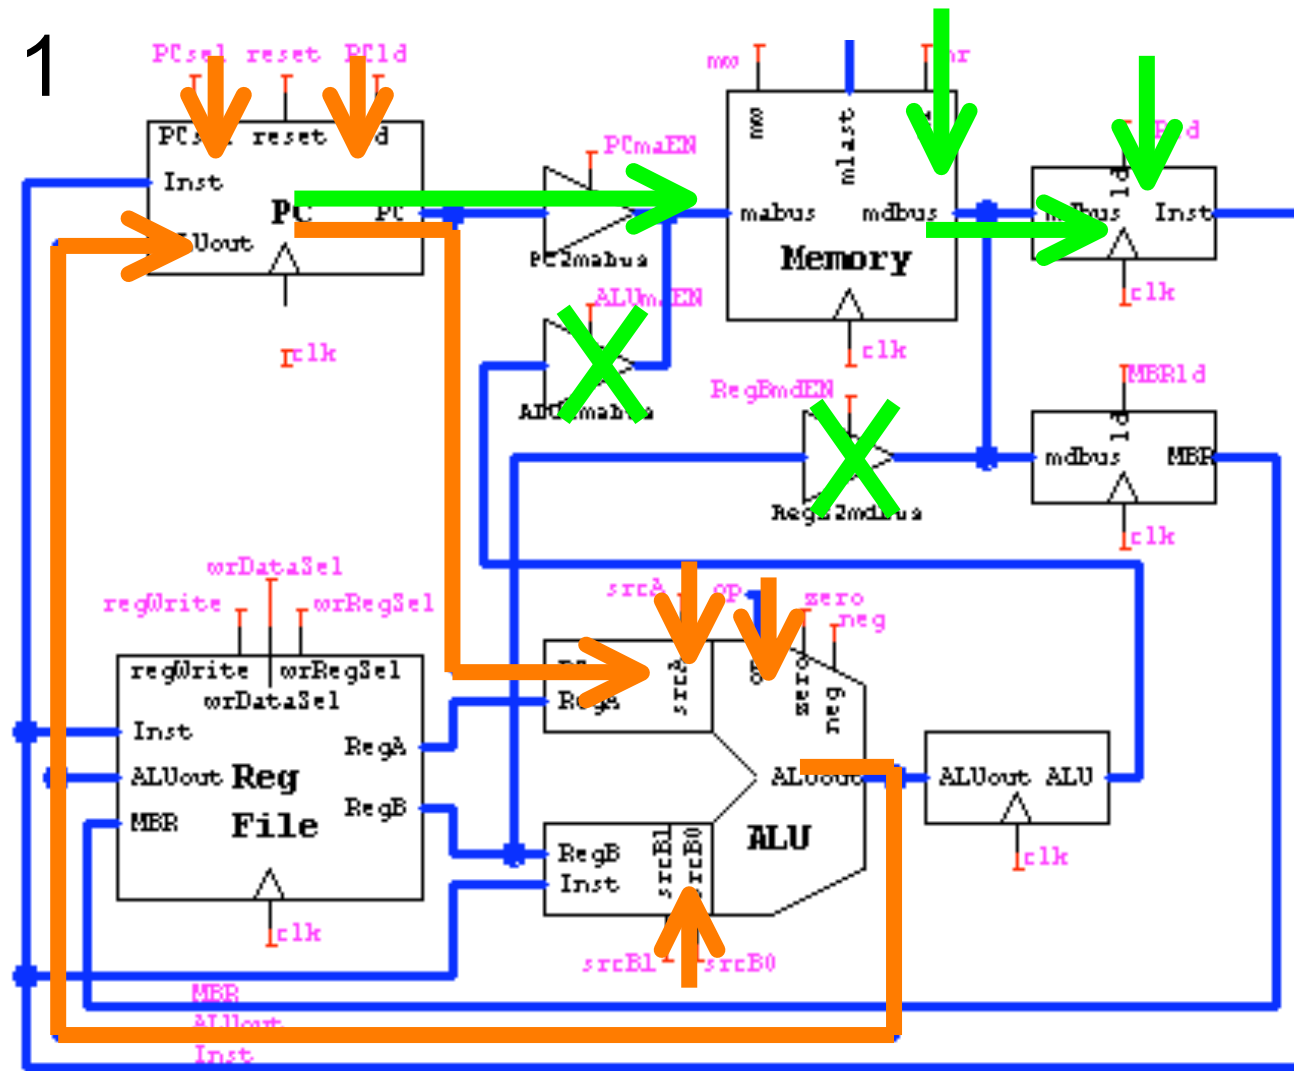
# RISC vs. CISC

---

- MIPS was one of the first RISC architectures. It was started about 20 years ago by [John Hennessy](#), one of the authors of our textbook.
- The architecture is similar to that of other RISC architectures, including [Sun's SPARC](#), [IBM](#) and [Motorola's PowerPC](#), and [ARM](#)-based processors.
- Older processors used complex instruction sets, or **CISC** architectures.
  - Many powerful instructions were supported, making the assembly language programmer's job much easier.
  - But this meant that the processor was more complex, which made the hardware designer's life harder.
- Many new processors use reduced instruction sets, or **RISC** architectures.
  - Only relatively simple instructions are available. But with high-level languages and compilers, the impact on programmers is minimal.
  - On the other hand, the hardware is much easier to design, optimize, and teach in classes.
- Even most current CISC processors, such as Intel 8086-based chips, are now implemented using a lot of RISC techniques.

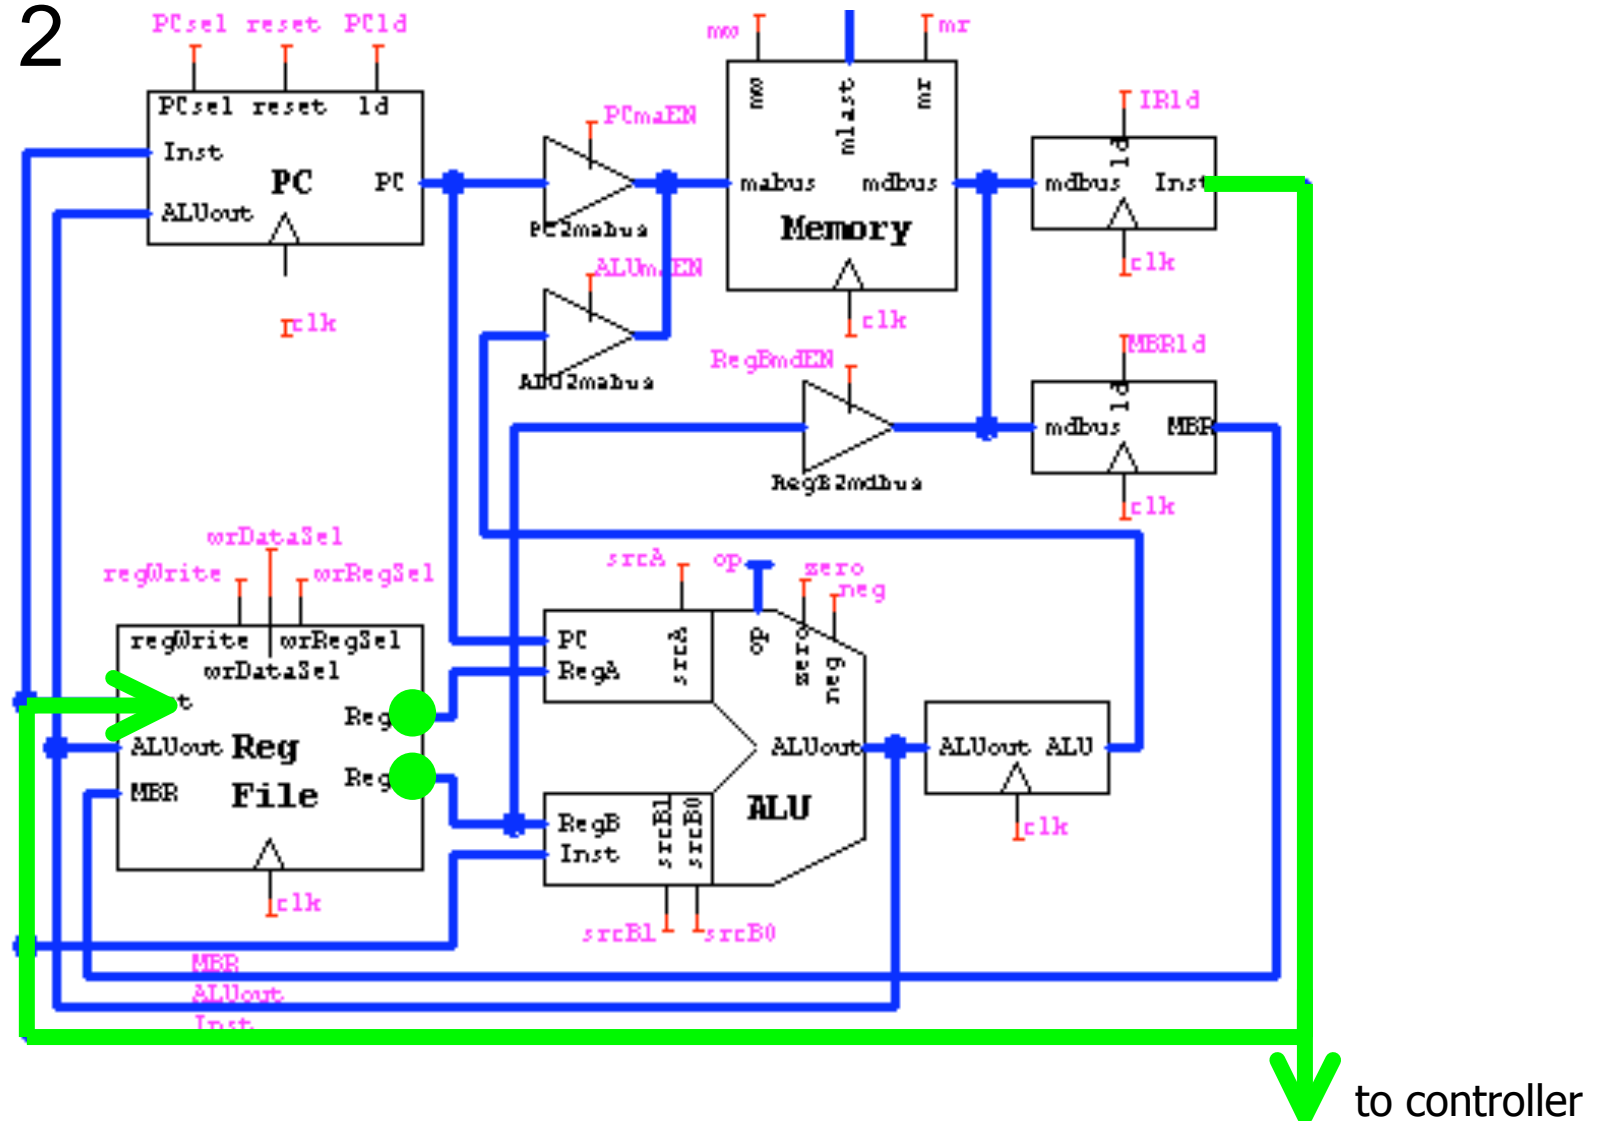
# Tracing an Instruction's Execution

- Step 1



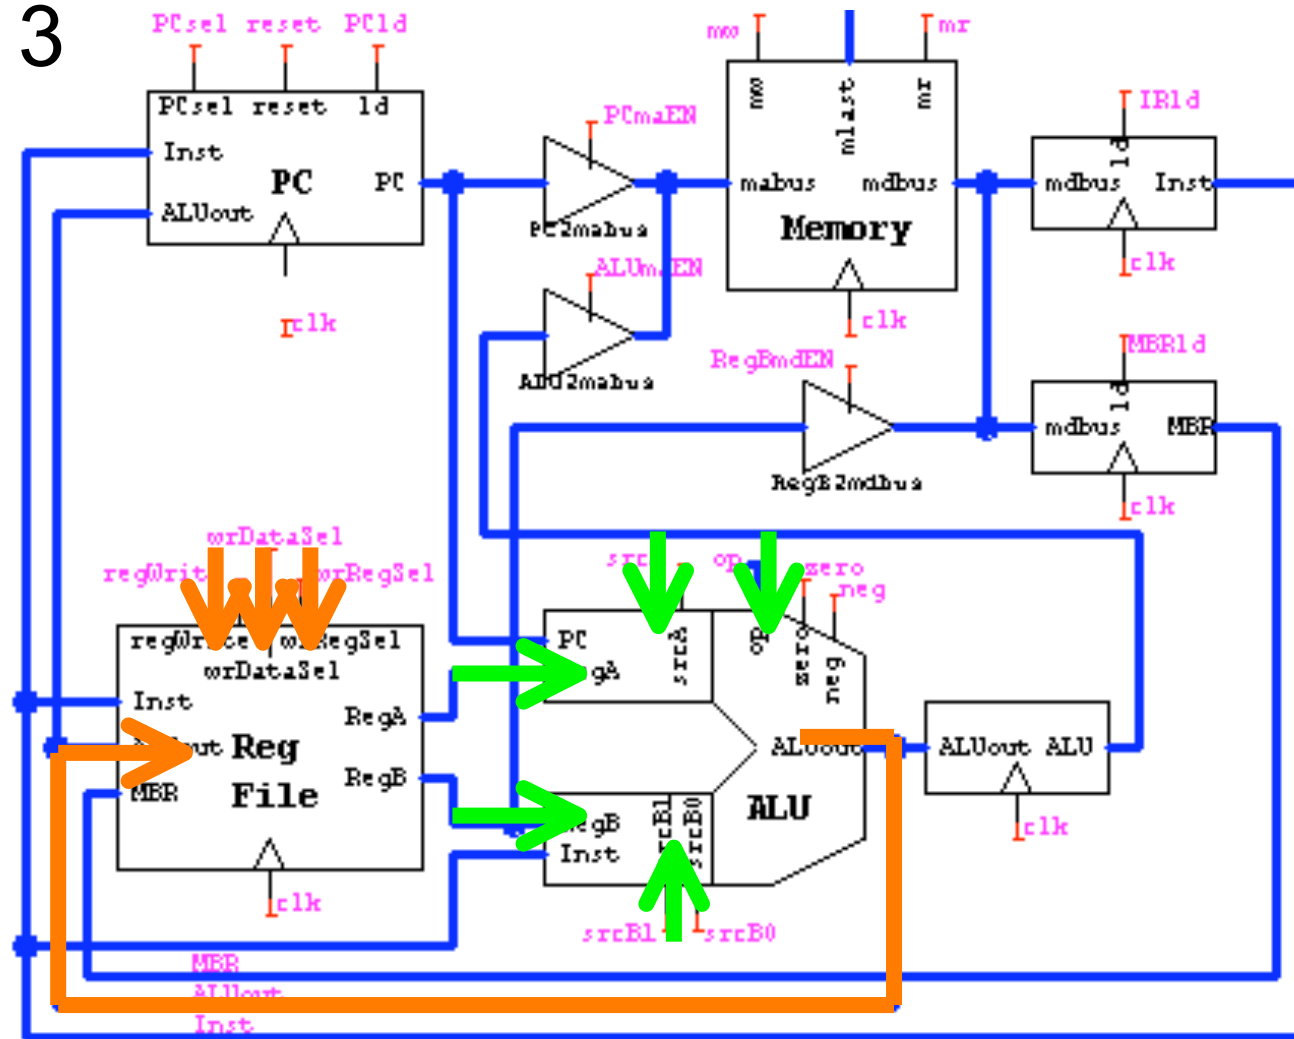
# Tracing an Instruction's Execution

- Step 2



# Tracing an Instruction's Execution

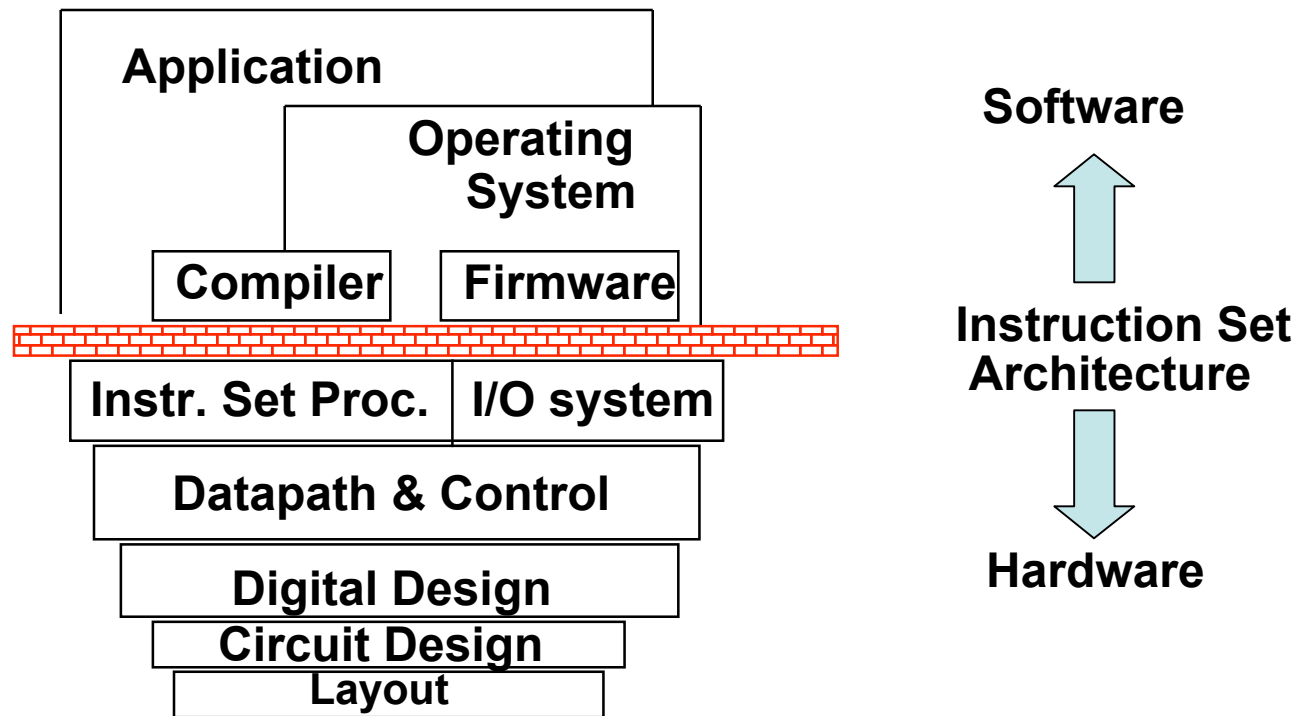
- Step 3





# Instruction Set Architecture

---




# Computer Organization

---

- Computer organization is the *implementation* of the machine, consisting of two components:
  - High-level organization
    - Memory system widths, bus structure, CPU internal organization, ...
  - Hardware
    - Precise logic and circuit design, mechanical packaging, ...
- Many implementations are possible for an ISA !!!
  - Intel i386, i486, Pentium, Pentium II, Core2Duo, QuadCore...
  - Performance, complexity, cost differences....
- Invisible to the programmer (mostly)!

# Moore's Law: 2x transistors every 18 months

---



# Technology => dramatic change

## □ Processor

- ➡ logic capacity: about 30% increase per year
- ➡ clock rate: about 20% increase per year

Higher logic density gave room for instruction pipeline & cache

## □ Memory

- ➡ DRAM capacity: about 60% increase per year (4x every 3 years)
- ➡ Memory speed: about 10% increase per year
- ➡ Cost per bit: about 25% improvement per year

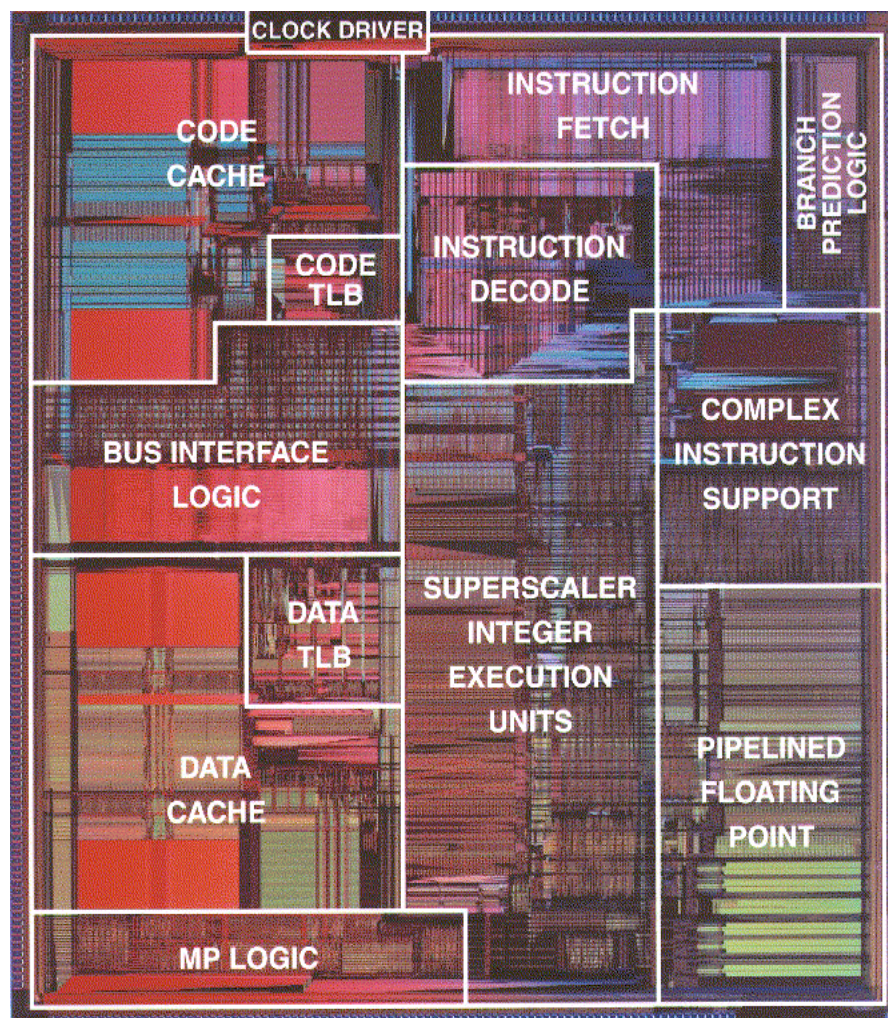
Performance optimization no longer implies smaller programs

## □ Disk

- ➡ Capacity: about 60% increase per year

Computers became lighter and more power efficient

# Pentium Die Photo



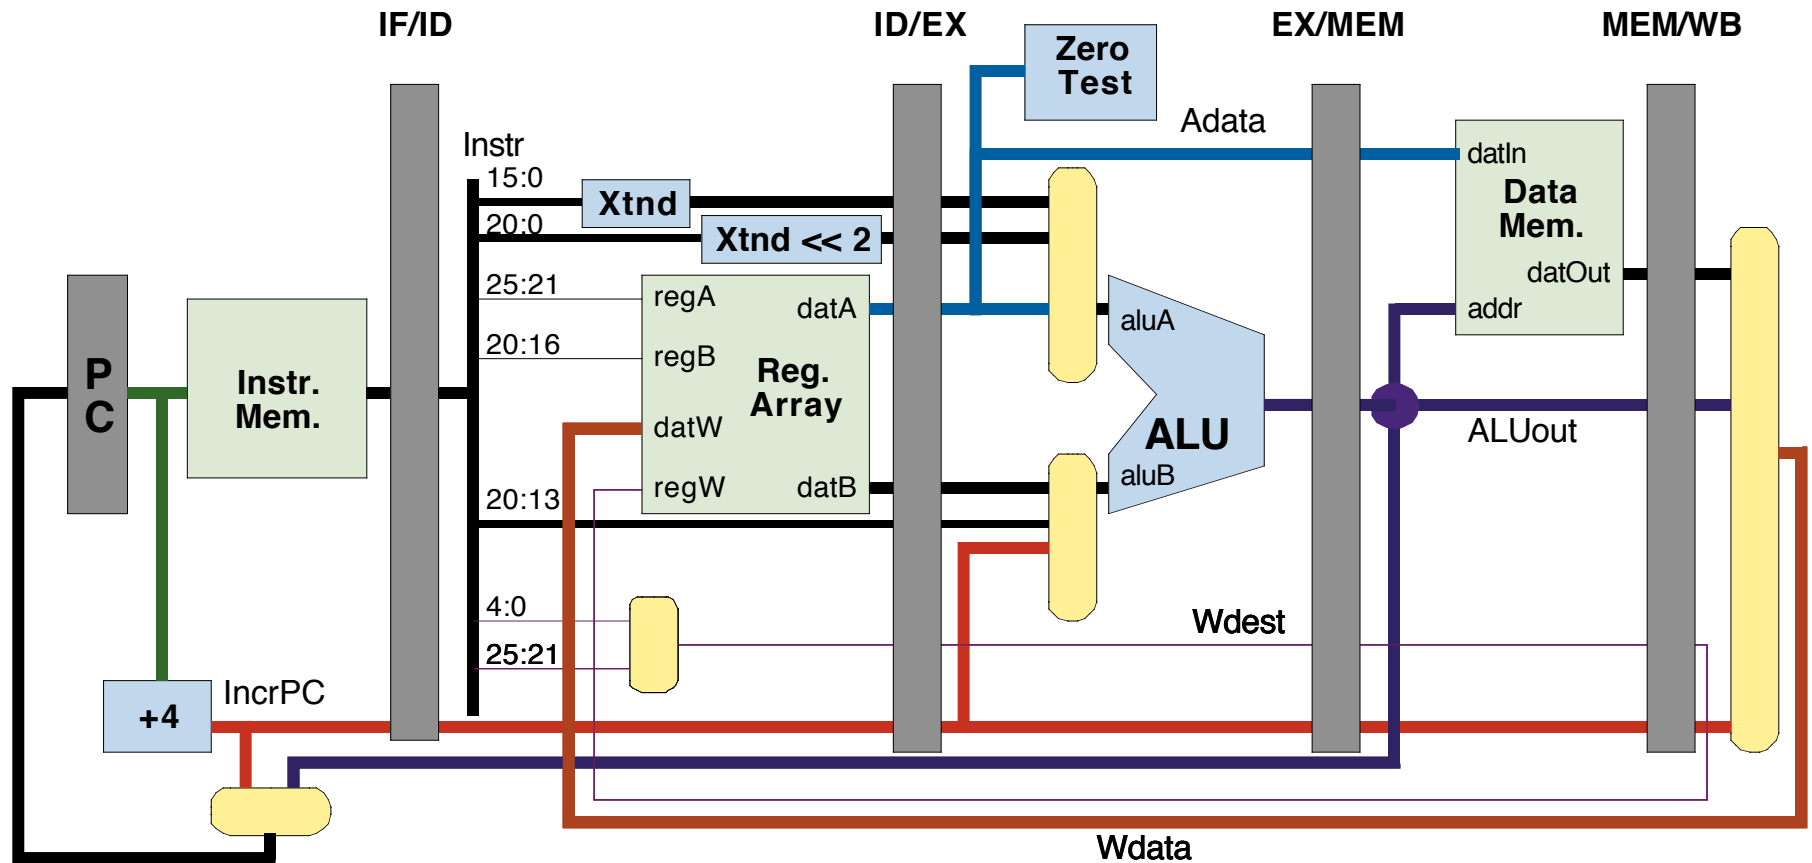
- 3,100,000 transistors
- 296 mm<sup>2</sup>
- 60 MHz
- Introduced in 1993
  - 1<sup>st</sup> superscalar implementation of IA32

# Four Organization Concepts/Trends

---

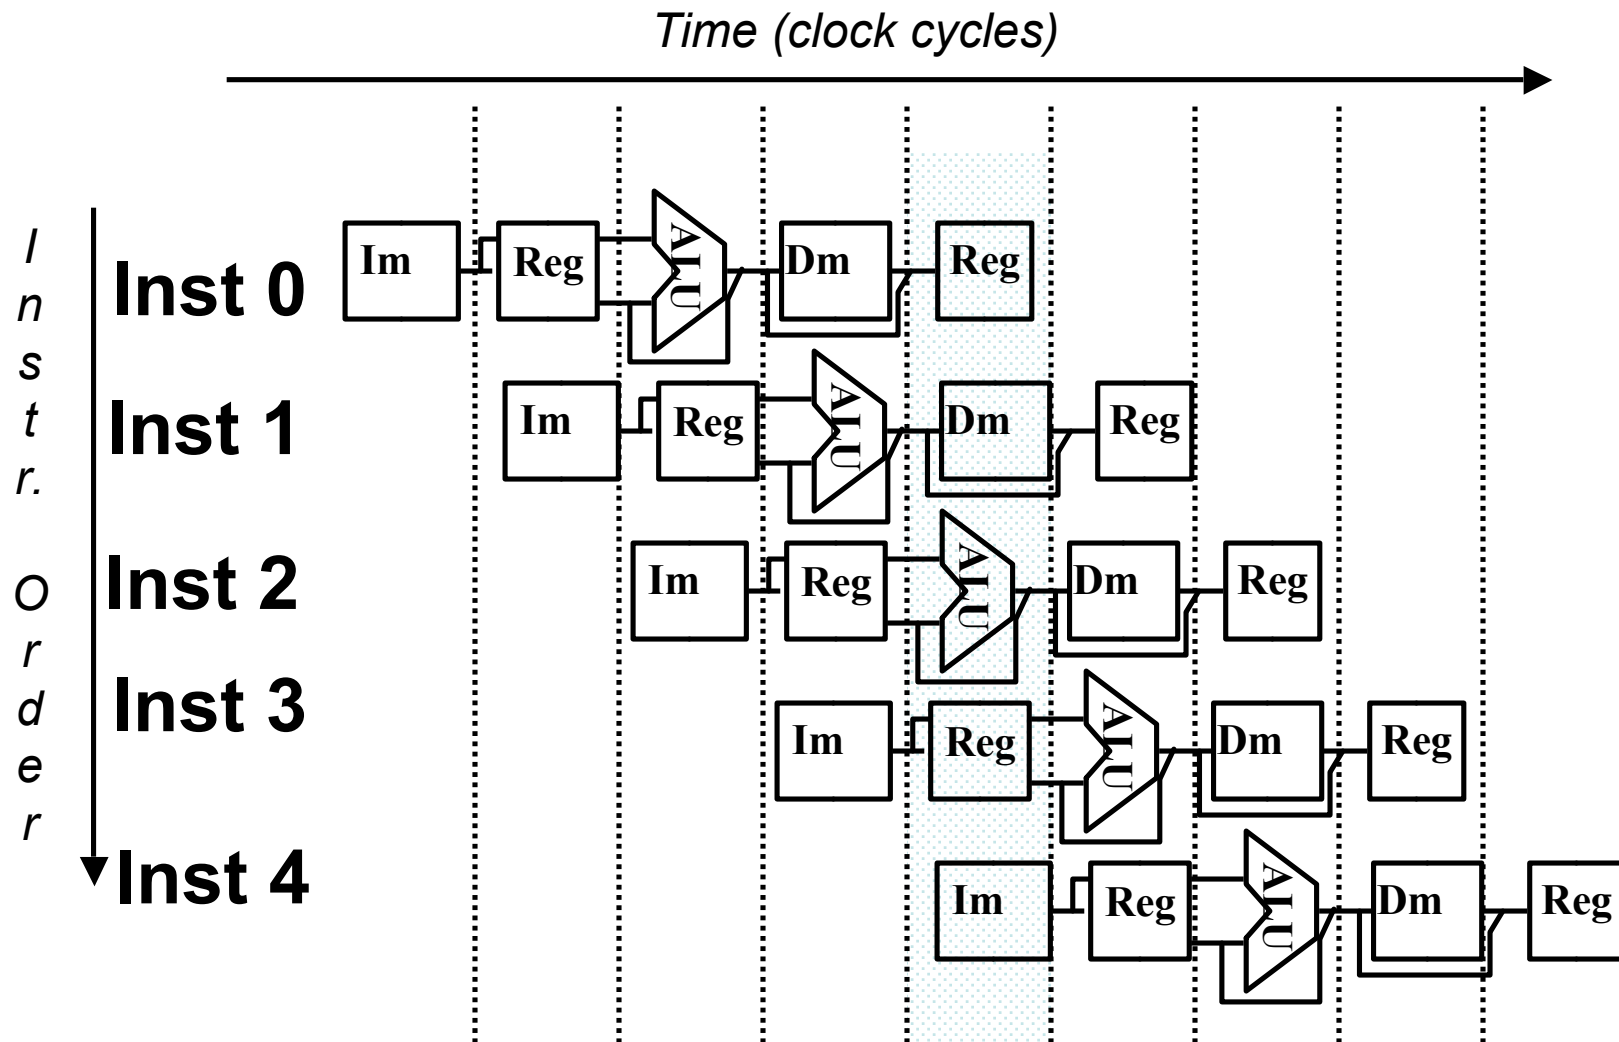
- Pipelining
  - Frequency
- Cache memory
  - To keep processor running must overcome memory latency
- Superscalar execution
  - Out of order execution
- Multi-core

# Pipelined Datapath



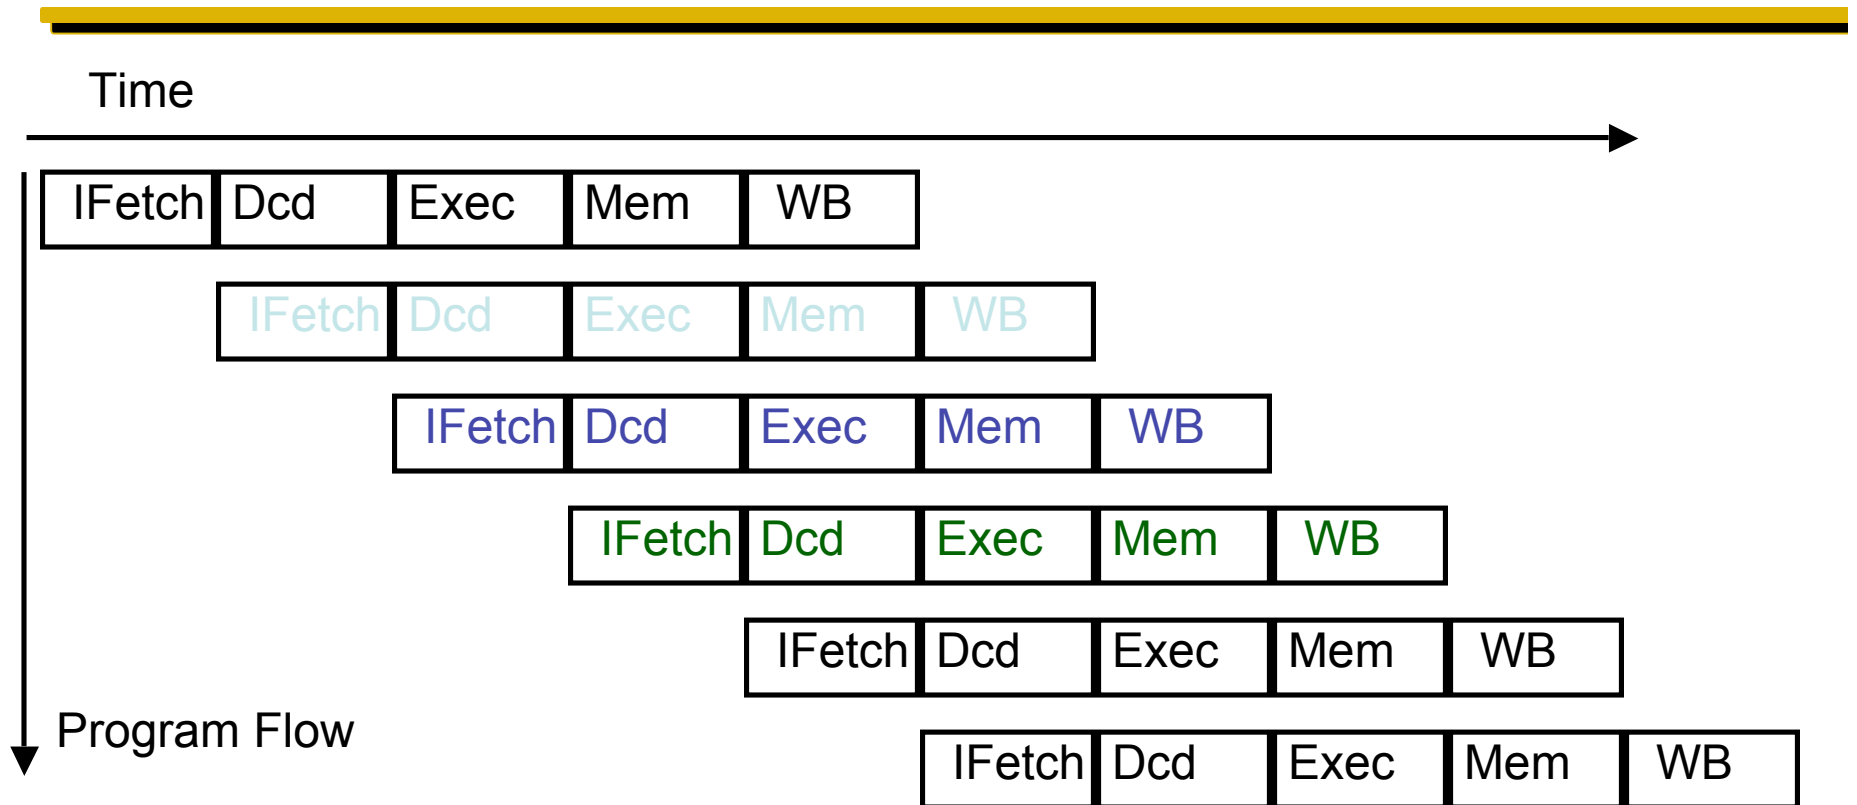
- Pipe Registers
  - Inserted between stages
  - Labeled by preceding & following stage

# Pipelining: Use Multiple Resources





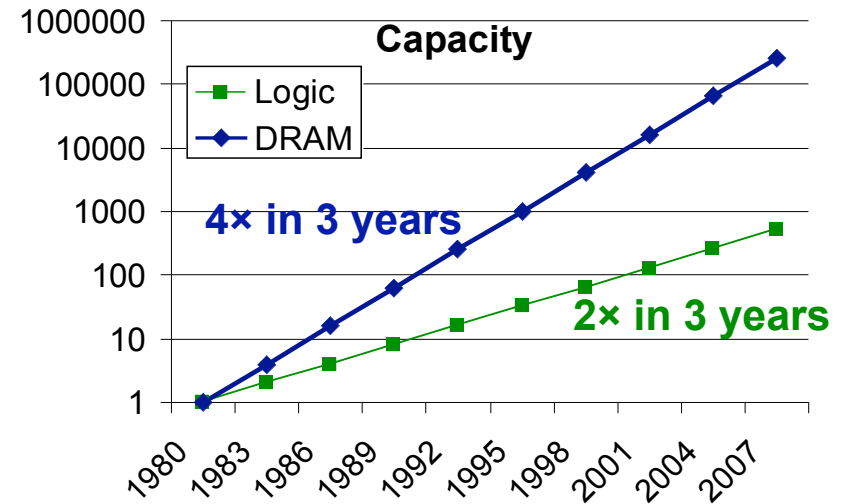

# Conventional Pipelined Execution Representation



- Improve performance by increasing instruction throughput


*Ideal speedup?  $N$  Stages and  $I$  instructions*

# Technology Trends and Performance



- Computing capacity: 4x per 3 years
  - If we could keep all the transistors busy all the time
  - Actual: 3.3x per 3 years
- Moore's Law: Performance is doubled every ~18 months
  - Trend is slowing: process scaling declines, power is up


# Processor-DRAM Memory Gap (latency)



# Cache Memory

---

- Small, fast storage used to improve average access time to slow memory.
- Exploits spatial and temporal locality
- In computer architecture, almost everything is a cache!
  - Registers a cache on variables
  - First-level cache a cache on second-level cache
  - Second-level cache a cache on memory
  - Memory a cache on disk (virtual memory)



# Modern Processor

This is an  
AMD Opteron CPU

Total Area: 193 mm<sup>2</sup>

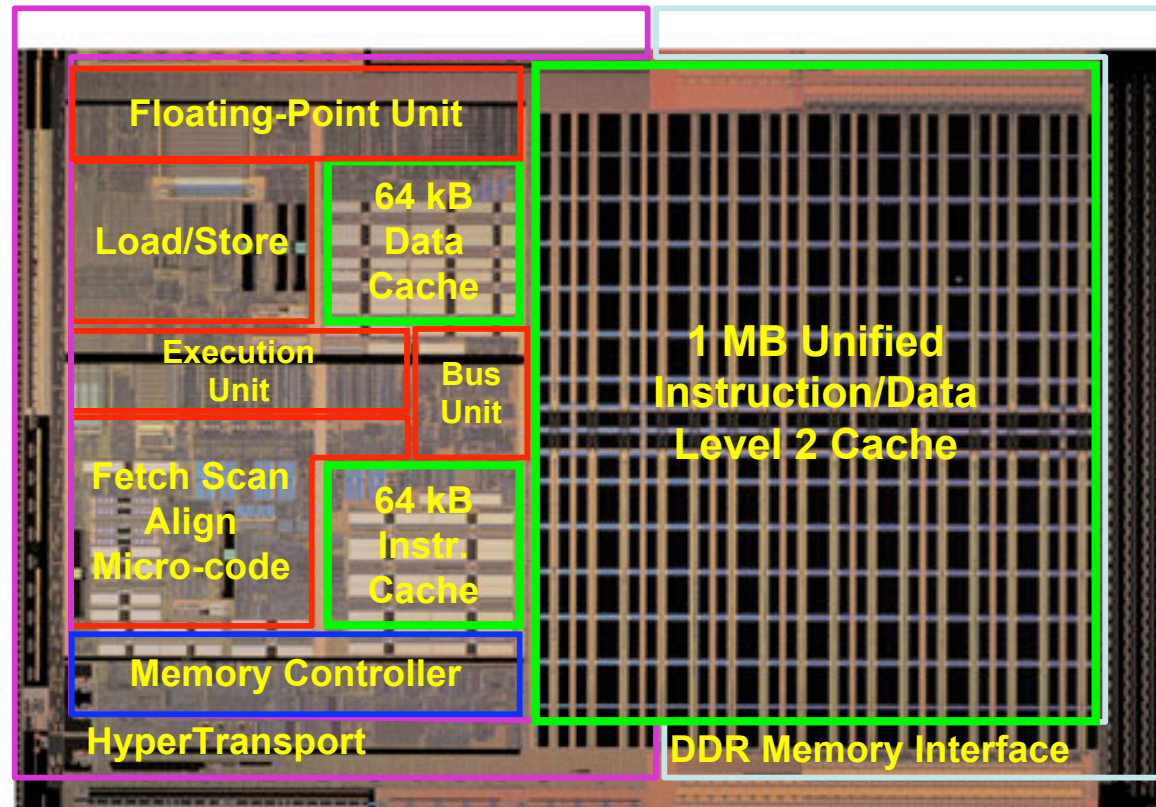
Look at the relative  
sizes of each block:

**50% cache**

**23% I/O**

**20% CPU logic**

**+ extra stuff**



CPU's dedicate > 50% area for **cache memory**.



# CPI – Cycles Per Instruction

---


- CPUs work according to a clock signal
  - Clock cycle is measured in nsec ( $10^{-9}$  of a second)
  - Clock frequency ( $= 1/\text{clock cycle}$ ) measured in GHz ( $10^9\text{cyc/sec}$ )
- Instruction Count (IC)
  - Total number of instructions executed in the program

$$\text{CPI} = \frac{\text{\#cycles required to execute the program}}{\text{IC}}$$

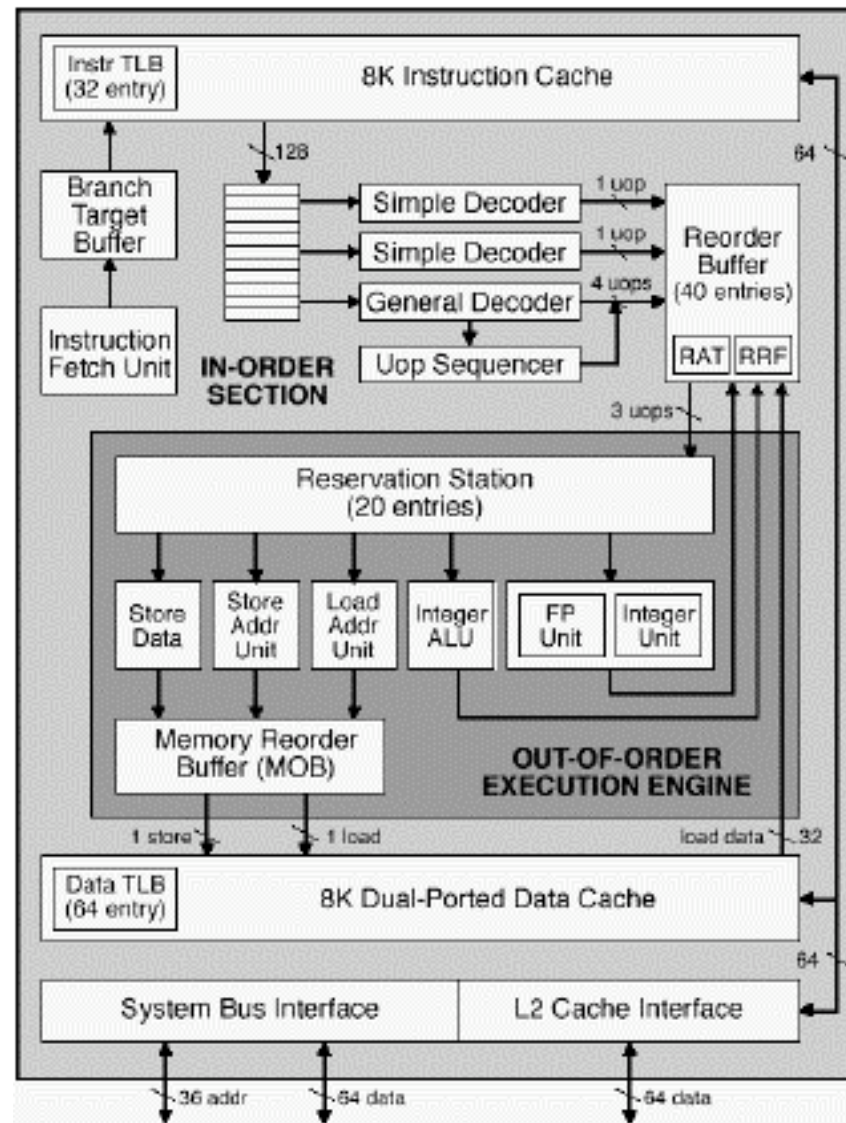
- CPI – Cycles Per Instruction
  - Average #cycles per Instruction (in a given program)
  - IPC ( $= 1/\text{CPI}$ ) : Instructions per cycles

# Relating time to system measures

- Iron Law: Performance = 1/execution time
- Suppose that for some program we have:
  - T seconds running time (the ultimate performance measure)
  - C clock ticks, I instructions, P seconds/tick (performance measures of interest to the system designer)
- $T \text{ secs} = C \text{ ticks} \times P \text{ secs/tick}$   
 $= (I \text{ inst}/I \text{ inst}) \times C \text{ ticks} \times P \text{ secs/tick}$
- $T \text{ secs} = I \text{ inst} \times (C \text{ ticks}/I \text{ inst}) \times P \text{ secs/tick}$



# Pentium 4 Block Diagram



Microprocessor Report



# Out Of Order Execution

---

- Look ahead in a window of instructions and find instructions that are *ready to execute*
  - Don't depend on data from previous instructions still not executed
  - Resources are available
- Out-of-order execution
  - Start instruction execution before execution of a previous instructions
- Advantages:
  - Help exploit Instruction Level Parallelism (ILP)
  - Help cover latencies (e.g., L1 data cache miss, divide)

# Data Flow Analysis

- Example:

(1)  $r1 \leftarrow r4 / r7$  ; assume divide takes 20 cycles

(2)  $r8 \leftarrow r1 + r2$


(3)  $r5 \leftarrow r5 + 1$

(4)  $r6 \leftarrow r6 - r3$


(5)  $r4 \leftarrow r5 + r6$

(6)  $r7 \leftarrow r8 * r4$


Data Flow Graph




**In-order execution**



**Out-of-order execution**



# OOOE – General Scheme




- Fetch & decode instructions in parallel but in order, to fill inst. pool
- Execute ready instructions from the instructions pool
  - All the data required for the instruction is ready
  - Execution resources are available
- Once an instruction is executed
  - signal all dependant instructions that data is ready
- Commit instructions in parallel but in-order
  - Can commit an instruction only after all preceding instructions (in program order) have committed

# Out Of Order Execution – Example

---

- Assume that executing a divide operation takes 20 cycles

(1)     **r1**  $\leftarrow$  r5 / r4  
(2)     **r3**  $\leftarrow$  **r1** + **r8**  
(3)     **r8**  $\leftarrow$  r5 + 1  
(4)     **r3**  $\leftarrow$  r7 - 2  
(5)     r6  $\leftarrow$  r6 + r7




- Inst2 has a REAL (flow) dependency on r1 with Inst1
  - It cannot be executed in parallel with Inst1
- Can successive instructions pass Inst2 ?
  - Inst3 cannot since Inst2 must read **r8** before Inst3 writes to it
  - Inst4 cannot since it must write to **r3** after Inst2
  - Inst5 can

# Overcoming False Dependencies

---

- OOOE creates new dependencies
  - **WAR (write after read)**: write to a register which is read by an earlier inst.
    - (1)  $r3 \leftarrow r2 + r1$
    - (2)  $r2 \leftarrow r4 + 3$
  - **WAW (write after write)**: write to a register which is written by an earlier inst.
    - (1)  $r3 \leftarrow r1 + r2$
    - (2)  $r3 \leftarrow r4 + 3$
- These are *false dependencies*
  - There is no missing data
  - Still prevent executing instructions out-of-order
- Solution: *Register Renaming*



# Pentium-4 Die Photo



1st Pentium4: 9.5 M transistors

- EBL/BBL - Bus logic, Front, Back
- MOB - Memory Order Buffer
- Packed FPU - MMX Fl. Pt. (SSE)
- IEU - Integer Execution Unit
- FAU - Fl. Pt. Arithmetic Unit
- MIU - Memory Interface Unit
- DCU - Data Cache Unit
- PMH - Page Miss Handler
- DTLB - Data TLB
- BAC - Branch Address Calculator
- RAT - Register Alias Table
- SIMD - Packed Fl. Pt.
- RS - Reservation Station
- BTB - Branch Target Buffer
- IFU - Instruction Fetch Unit (+I\$)
- ID - Instruction Decode
- ROB - Reorder Buffer
- MS - Micro-instruction Sequencer

# An Exciting Time for CE - CMPs





# Computer System Structure

