



Aggregate Functions & Nested Queries

Dr. Munesh Singh

Indian Institute of Information Technology
Design and Manufacturing,
Kancheepuram
Chennai-600127

February 15, 2019





Aggregate functions

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- These functions operate on the multiset of values of a column of a relation, and return a value
 - 1 avg: average value
 - 2 min: minimum value
 - 3 max: maximum value
 - 4 sum: sum of values
 - 5 count: number of values



Aggregate functions

Aggregate Functions & Nested Queries

Dr. Munesh Singh

- Find the average salary of instructors in the computer science department
`select avg(salary) from instructor where deptname='cse';`
- Find the total number of the instructor who teach a course in Spring 2018 semester
`select count(distinct ID) from teaches where semester='Spring' and year='2010';`
- Find the number of tuple in the course relation
`select count(*) from course;`



Aggregate functions-Group By and having

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Find the salary of instructors in each department
`select deptname, avg(salary) as avg_salary from instructor group by deptname;`
- Find the names and average salaries of all departments whose salary is greater than 42000
`select deptname, avg(salary) from instructor group by deptname having avg(salary)>42000;`
- **Note:** predicates in the having clause are applied after the formation of groups whereas predicates in the where clause are applied before forming groups



Null Values and Aggregates

Aggregate Functions & Nested Queries

Dr. Munesh Singh

- Total all salaries
`select sum(salary) from instructor;`
- Above statement ignores null amounts
- Result is null if there is no non-null amount
- All aggregate operations except count(*) ignore tuples with null values on the aggregated attributes
- what if collection has only null values?
 - count returns 0
 - all other aggregates return null



Nested Sub-Queries

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- SQL provides a mechanism for the nesting of subqueries
- A subquery is a **select-from-where** expression that is nested within another query
- The nesting can be done in the following SQL query

```
select  $A_1, A_2, \dots, A_n$   
from  $r_1, r_2, \dots, r_n$   
where P
```
- as follows:
- A_i can be replaced by a subquery that generate a single value
- r_i can be replaced by any valid subquery
- P can be replaced with an expression of the form:
 - $B < operation > (subquery)$
 - where B is an attribute and $< operation >$ to be defined later



Sub-Queries in the Where Clause

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- A common use of subqueries is to perform tests:
 - For set membership
 - For set comparisons
 - For set cardinality



Set Membership

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Find courses offered in Fall 2009 and in Spring 2010
select distinct courseid from section where semester='Fall'
and year=2009 and courseid in (select courseid from
section where semester='Spring' and year=2010)
- Find courses offered in Fall 2009 and not in Spring 2010
select distinct courseid from section where semester='Fall'
and year=2009 and courseid not in (select courseid from
section where semester='Spring' and year=2010)





Set Membership Cont...

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Find the total number of (distinct) students who have taken course sections taught by the instructor with ID 10101

```
select count(distinct ID) from  
takes(courseid,secid,semester,year) in (select courseid  
secid,semester,year from teaches where teaches.ID=10101)
```



Set Comparison - "some" Clause

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Find the names of instructor with salary greater than that of some (at least one) instructor in the Biology department
`select distinct T.name from instructor as T, instructor as S
where T.salary > S.salary and S.deptname='Biology';`
- Same query using `> some` clause
`select name from instructor where salary > some (select
salary from instructor where dept name='Biology');`



Set Comparison - "all" Clause

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Find the names of instructor whose salary is greater than the salary of all instructor in Biology department
`select name from instructor where salary > all (select salary from instructor where dept name='Biology');`



Test for Empty Relations

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- There **exists** construct return the value true if the argument subquery is nonempty
- $\text{exist } r \Leftrightarrow r \neq \Phi$
- $\text{not exist } r \Leftrightarrow r = \Phi$
- Find courses taught in both Fall 2009 and in Spring 2010 semester

```
select courseid from section as S where semester='Fall'  
and year=2009 and exists (select * from section as T  
where semester='Spring' and year=2010 and  
S.courseid=T.courseid);
```



Use of "not exists" Clause

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Find all students who have taken all courses offered in the Biology department.
`select distinct S.ID,S.name from student as S where not exists((select courseid from course where deptname='Biology') except (select T.courseid from takes as T where S.ID=T.ID))`
- First nested query list all courses offered in Biology
- Second nested query lists all courses a particular student took



Test for Absence of Duplicate Tuples

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- The unique construct tests whether a subquery has any duplicate tuple in its result
- The unique construct evaluates to "true" if a given subquery contains no duplicates
- Find all courses that were offered at most once in 2009

```
select T.courseid from course as T where unique(select  
R.courseid from section as R where T.courseid=R.courseid  
and R.year=2009);
```



Subqueries in the From Clause

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- SQL allows a subquery expression to be used in the **from** clause
- Find the average instructor salaries of those departments where the average salary is greater than 42000.
`select deptname,avgs from (select deptname, avg(salary) as avgs from instructor group by deptname) where avgs>42000;`
- same query in another way
`select deptname,avgs from (select deptname, avg(salary) as from instructor group by deptname) as deptavg(dname,avgs) where avgs>42000;`



With Clause

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs
- Find all department with the maximum budget
with maxbudget(value) as(select max(budget) from department) select department.name from department, maxbudget where department.budget=maxbudget.value



Complex Queries using With Clause

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Find all department where the total salary is greater than the average of the total salary at all departments
with depttotal(deptname,value) as (select deptname, sum(salary) from instructor group by deptname), deptavg(value) as(select avg(value) from depttotal) select deptname from depttotal,deptavg where depttotal.value > deptavg.value;



Subqueries in the Select Clause

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

Scalar Subquery

- Scalar subquery is one which is used where a single value is expected
- List all department along with the number of instructors in each department

```
select deptname,(select count(*) from instructor  
where department.deptname=instructor.deptname)  
as numinstructors from department
```



Modification of the Database

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Deletion of tuples from a given relation
- Inserting of new tuples into a given relation
- Updating of values in some tuples in a given relation
- delete all instructors
delete from instructor
- delete all instructor from Finance department
delete from instructor where deptname='Finance'



Modification of the Database

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Delete all tuples on the instructor relation for those instructors associated with a department located in the Watson building
`delete from instructor where deptname in (select deptname from department where building='Watson');`



Modification of the Database

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Delete all instructors whose salary is less than the average salary of instructors
`delete from instructor where salary < (select avg(salary) from instructor);`
- **Problem:** as we delete tuples from deposit, the average salary changes
- **Solution used in SQL**
 - First, compute `avg(salary)` and find all tuples to delete
 - Next, delete all tuples found above (without recomputing `avg` or retesting the tuples)



Insertion

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Add a new tuple to course
**insert into course values('cs-117','Database System',
'CSE',4);**
- or equivalently
**insert into course (courseid,title,deptname,credit)
values('cs-117','Database System', 'CSE',4);**
- Add a new tuple to student with totalcredit set to null
**insert into student values('3003','Green',
'Finance',0);**



Insertion (Cont...)

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Add all instructors to the student relation with totalcredit set to 0
insert into student select ID, name,deptname,0 from instructor
- The select from where statement is evaluated fully before any of its results are inserted into the relation



Updates

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Increase salaries of instructors whose salary is over 1000000 by 3% and all others by a 5%
- Write two update statements:
update instructor set salary=salary*1.03 where salary > 100000;
update instructor set salary=salary*1.05 where salary <= 100000;
- The order is important
- Can be done better using the case statement



Case Statement for Conditional Updates

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Same query as before but with case statement
update instructor set salary=case
when salary <= 100000 then salary*1.05
else salary*1.03
end



Views

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- In some cases, it is not desirable for all users to see the entire logical model (that is, all actual stored in the database.)
- Consider a person who needs to know an instructors name and department, but not the salary. This person should see a relation describe, in SQL by
`select ID, name, deptname from instructor`
- A view provides a mechanism to hide certain data from the view of certain users
- relation that is not of the conceptual model but is made available to a user as a "virtual relation" is called a view



Views Definition

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- A view is defined using the create view statement which has the form
create view v as *< query expression >*
- where *< queryexpression >* is any legal SQL expression
- The view name is represented by v
- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates
- View definition is not the same as creating a new relation by evaluating the query expression
 - Rather, a view definition causes the saving of an expression; the expression is substituted into queries using the view



Example Views

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- A view of instructors without their salary
create view faculty as select id,name,deptname from instructor
- Find all instructors in Biology department **select name from faculty where deptname='Biology'**
- Create a view of department salary totals **create view departmenttotalsalary(deptname, totalsalary) as select deptname, sum(salary) from instructor group by deptname**



Views Defined Using Other Views

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- **create view physics_fall_2009_watson as select
courseid,roomno fromselect
course.courseid,secid,building,roomno
from (select course.courseid,secid,building,roomno
course,section
where course.courseid=section.courseid
and course.deptname='Physics'
and section.semester='Fall'
and section.year='2009');**



Views Defined Using Other Views

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- One view may be used in the expression defining another view
- A view relation V_1 is said to **depend directly** on a view relation V_2 if V_2 is used in the expression defining V_1
- A view relation V_1 is said to **depend on** view relation V_2 if either V_1 depends directly to V_2 or there is a path of dependencies from V_1 to V_2
- A view relation v is said to be recursive if it depends on itself



Views Expansion

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- A way to define the meaning of views defined in terms of other views
- Let view V_1 be defined by an expression e_1 that may itself contain uses of view relations
- View expansion of an expression repeats the following replacement step:
repeat
Find any view V_i in e_1
Replace the view relation V_j by the expression defining V_j
until no more view relations are present in e_1
- As long as the view definitions are not recursive, this loop will terminate



Recursive View

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- In SQL, recursive queries are typically built using these components:
 - ① A non-recursive seed statement
 - ② A recursive statement
 - ③ A connection operator
The only valid set connection operator in a recursive view definition is UNION ALL
 - ④ A terminal condition to prevent infinite recursion
- View expansion of an expression repeats the following replacement step:
repeat
Find any view V_i in e_1
Replace the view relation V_j by the expression defining V_j
until no more view relations are present in e_1
- As long as the view definitions are not recursive, this loop will terminate



Recursive View-Example

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- In the context of a relation flights:
- Find all the destinations that can be reached from 'Paris'

```
create table flights (  
  source varchar(40),  
  destination varchar(40),  
  carrier varchar(40),  
  cost decimal(5,0));
```

source	destination	carrier	cost
Paris	Detroit	KLM	7
Paris	New York	KLM	6
Paris	Boston	American Airlines	8
New York	Chicago	American Airlines	2
Boston	Chicago	American Airlines	6
Detroit	San Jose	American Airlines	4
Chicago	San Jose	American Airlines	2



Recursive View-Example

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- **create recursive view**
reachable_from(source,destination,depth) as (select
root.source,root.destination, 0 as depth
from flight as root
where root.source='Paris'
union all
select in1.source, out1.destination,in1.depth+1 from
reachable_from as in1,flights as out1
where in1.destination=out1.source and
in1.depth<= 100);



Recursive View-Example

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

Result of the previous query

- Get the result by simple selection on the view:
`select distinct source,destination from reachable_from`

source	destination	carrier	cost
Paris	Detroit	Jet	7
Paris	New York	Jet	6
Paris	Boston	indigo	8
New York	Chicago	indigo	2
Boston	Chicago	indigo	6
Detroit	San Jose	indigo	4
Chicago	Detroit	indigo	7



The Power of Recursion

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Recursive View makes it possible to write queries, such as transitive closure queries, that can not be written without recursion or iteration
- Intuition: Without recursion, a non-recursive non interactive program can perform only a fixed number of joins of flights with itself
 - This can give only a fixed number of levels of reachable destinations
 - Given a fixed non-recursive query, we can construct a database with a greater number of levels of reachable destinations on which the query will not work



The Power of Recursion

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Computing transitive closure using iteration, adding successive tuples to `reachable_from`
 - The next slide shows a `flights` relation
 - Each step of the iterative process constructs an extended version of `reachable_from` from its recursive definition
 - The final result is called the fixed point of the recursive view definition.
- Recursive view are required to be monotonic. That is, if we add tuples to `flights` the view `reachable_from` contains all of the tuples it contained before, plus possible more



The Power of Recursion

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

source	destination	carrier	cost
Paris	Detroit	Jet	7
Paris	New York	Jet	6
Paris	Boston	indigo	8
New York	Chicago	indigo	2
Boston	Chicago	indigo	6
Detroit	San Jose	indigo	4
Chicago	Detroit	indigo	7

Iteration#	Tuples in closure
0	Detroit, New York, Boston
1	Detroit, New York, Boston, San Jose, Chicago
2	Detroit, New York, Boston, San Jose, Chicago



Update of a View

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Add a new tuple to faculty view which we defined earlier
insert into faculty values('30765','Green','Music');
This insertion must be represented by the insertion of the
tuple ('30765','Green','Music',null)
into the instructor relation



Some Updates cannot be Translated Uniquely

- **create view instructor_info as**
select ID, name, building from instructor, department
where instructor.dept_name=department.dept_name;
- **insert into instructor_info**
values('69987','White','Taylor');
 - which department, if multiple department in Taylor?
 - what if no department is in Taylor?
- Most SQL implementations allow updates only on simple views
 - The **form** clause has only one database relation
 - The **select** clause contains only attribute names of the relation, and does not have any expressions, aggregates, or **distinct** specification
 - Any attribute not listed in the **select** clause can be set to null
 - The query does not have a **group** by or **having** clause



Authorization

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Forms of authorization on parts of the database:
 - **Read**- allows reading, but not modification of data
 - **Insert**- allows insertion of new data, but not modification of existing data
 - **Update**- allows modification, but not deletion of data
 - **Delete**- allows deletion of data
- Forms of authorization to modify the database schema
 - **Index**- allows creation and deletion of indices
 - **Resources**- allows creation of new relations
 - **Alteration**- allows addition or deletion of attributes in a relations
 - **Drop**- allows deletion of relations



Authorization

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- The grant statement is used to confer authorization:
 - **grant** *< privilegelist >*
on *< relationnameorviewname >* **to** *< userlist >*
- *< userlist >* is:
 - **grant** a user-id
 - **public**, which allows all valid users the privilege granted
 - A role (more on this later)
- Granting a privilege on a view does not imply granting any privileges on the underlying relations
- The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator)



Some Updates cannot be Translated Uniquely

- **select:** allows read access to relation, or the ability to query using the view
- Example: grant users U_1, U_2, U_3 select authorization on the instructor relation:
grant select on instructor to U_1, U_2, U_3
- **insert:** the ability to insert tuples
- **update:** the ability to update using the SQL update statement
- **delete:** the ability to delete tuples.
- **all privileges:** used as a short from for all the allowable privileges



Revoking Authorization in SQL

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- The **revoke** statement is used to revoke authorization
revoke < *privilegelist* >
on < *relationnameorviewname* > **from** < *userlist* >
- Example:
revoke select on branch **from** U_1, U_2, U_3
- < *privilege – list* > may be all to revoke all privileges the revokee may hold
- if < *revokee – list* > includes **public**, all users lose the privilege except those granted it explicitly
- if the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the revocation
- All privileges that depends on the privilege are being revoke are also revoked



Roles

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- **create role** instructor;
grant instructor to Amit;
- Privileges can be granted to roles:
grant select on takes **to** instructor;
- Roles can be granted to users, as well as to other roles
create role teaching_assistant
grant teaching_assistant to instructor;
- Instructor inherits all privileges to teaching_assistant
- Chain of roles
 - **create role** dean;
 - **grant** instructor **to** dean;
 - **grant** dean to ram;

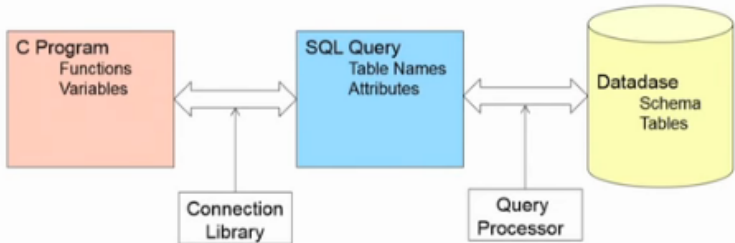


Accessing SQL From A Programming Language

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

Native Language < -- > Query Language





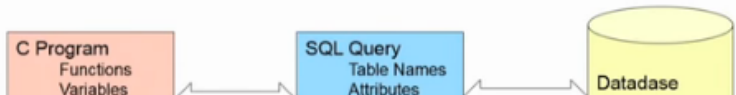
Accessing SQL From A Programming Language

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

Native Language < — > Query Language

- API (application programming interface) for a program to interact with a database server
- Application makes calls to
 - connect with the database server
 - Send SQL commands to the database server
 - Fetch tuples of result one-by-one into program variables
- Various Tools:
 - JDBC (Java Database Connectivity) work with java
 - ODBC (open Database Connectivity) works with C, C++, C#, Python.
 - Embedded SQL





JDBC

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

Native Language < -- > Query Language

- JDBC is a java API for communicating with database systems supporting SQL
- JDBC supports a variety of features for querying and updating data, and for retrieving results.
- JDBC also supports metadata retrieval, such as querying about relations present in the database and the names and types of relation attributes
- Model for communicating with database:
 - Open a connection
 - Create a "statement" object
 - Execute queries using the statement object to send queries and fetch results
 - Exception mechanism to handle error



ODBC

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

Native Language < — — > Query Language

- Open Database Connectivity (ODBC) standard
 - standard for application program to communicate with a database server
 - application program interface (API) to
 - open a connection with a database,
 - send queries and updates,
 - get back results



Application Program Connection With ODBC

Native Language < -- > Query Language

Application Program

(MySQLdb is a thin Python wrapper around _mysql.

It is compatible with the Python DB API)

```
#!/usr/bin/python
```

```
# -*- coding: utf-8 -*-
```

```
import MySQLdb as mdb
```

```
import sys
```

```
try:
```

```
    con = mdb.connect('localhost', 'testuser', 'test623', 'testdb')
```

```
    cur = con.cursor()
```

```
    cur.execute("SELECT VERSION()")
```

```
    ver = cur.fetchone()
```

```
    print "Database version : %s " % ver
```

```
except mdb.Error, e:
```

```
    print "Error %d: %s" % (e.args[0], e.args[1])
```

```
    sys.exit(1)
```

```
finally:
```

```
    if con:
```

```
        con.close()
```

Database Query

```
mysql> CREATE DATABASE testdb;
```

```
mysql> CREATE USER 'testuser'@'localhost' IDENTIFIED BY 'test623';
```

```
mysql> USE testdb;
```

```
mysql> GRANT ALL ON testdb.* TO 'testuser'@'localhost';
```

```
mysql> quit;
```

```
cur = con.cursor()
```

```
cur.execute("DROP TABLE IF EXISTS Writers")
```

```
cur.execute("CREATE TABLE Writers(Name VARCHAR(25))")
```

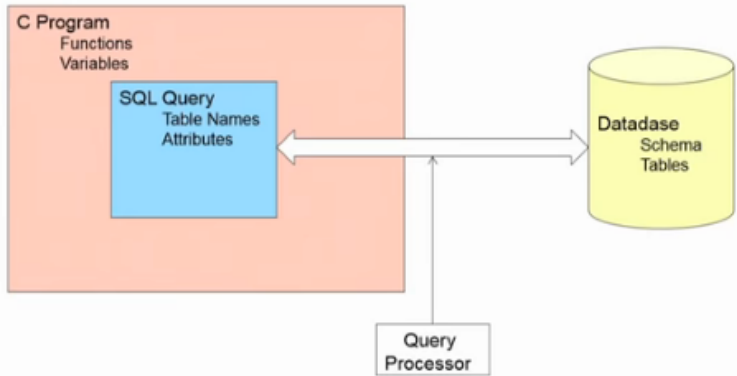
```
cur.execute("INSERT INTO Writers(Name) VALUES('Jack London')")
```

```
cur.execute("INSERT INTO Writers(Name) VALUES('Honore de Balzac')")
```



SQL Part of Application Programming Language(Embedded SQL)

Native Language < -- > Query Language



Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh



Triggers)

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

Native Language < -- > Query Language

- A **trigger** is a statement that is executed automatically by the system as a side effect of a modification to the database
- To design a trigger mechanism, we must:
 - specify the condition under the trigger is to be executed.
 - specify the actions to be taken when trigger executes
- Trigger introduced to SQL standard in:1999, but support even earlier using non-standard syntax by most databases



Triggering Events and Actions in SQL)

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

Native Language < -- > Query Language

- Triggering events can be **insert, delete, updates**
item Triggers on update can be restricted to specific attributes
 - For example, **after update of** takes on grade
- Values of attributes before and after an update can be referenced
 - **referencing old row as:** for deletes and updates
 - **referencing new row as:** for inserts and updates
- Trigger can be activated before an event, which can serve as extra constraints. For example, convert blank grades to null.



Trigger Example

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- **create trigger setnull_trigger before update of takes
referencing new row as nrow
for each row
where (nrow.grade=”)
begin atomic
set nrow.grade=null
end;**



When NoT To Use Triggers

- Triggers were used earlier for tasks such as
 - Maintaining summary data (e.g. total salary of each department)
 - Replicating databases by recording changes to special relations (called change or delta relations) and having a separate process that applies the changes over to a replica
- There are better ways of doing these now:
 - Databases today provide built in materialized view facilities to maintain summary data
 - Databases provide built-in support for replication
- Encapsulation facilities can be used instead of triggers in many cases
 - Define methods to update fields
 - Carry out actions as part of the update methods instead of through a trigger



When NoT To Use Triggers

Aggregate
Functions &
Nested
Queries

Dr. Munesh
Singh

- Risk of unintended execution of trigger, for example ,
when
 - Loading data from a backup copy
 - Replicating updates at a remote site
 - Trigger execution can be disabled before such actions