



# Transaction Management/ACID in DBMS

Dr. Munesh Singh

Indian Institute of Information Technology  
Design and Manufacturing,  
Kancheepuram  
Chennai-600127

March 25, 2019





# Transaction Management

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

## How transaction happens?

- A transaction is a program unit whose execution may change the contents of a database
- **Database before transaction**
- **Database after transaction**
- Transaction is used to represent a logical unit of database processing that must be completed in its entirety to ensure correctness
- Transaction is a collection of **Read, Update, Write**.



# ACID Properties of Transaction

## Atomicity

- **Atomicity (ALL or NONE)** Ensure that a transaction will run to completion as an indivisible unit. At the end no change occur to database or database change should be consistent manner.
- Example: Given  $A=2000$ ,  $B=3000$ ,  $\text{initial}=5000$   
 $\text{Rs}=500/-$  from  $A \rightarrow B$   
T1  
 $R(A,a)$   
 $a=500$   
 $W(A,a) < -$



# ACID Properties of Transaction

## Atomicity

- **Atomicity (ALL or NONE)** Ensure that a transaction will run to completion as an indivisible unit. At the end no change occur to database or database change should be consistent manner.
- Example: Given  $A=2000$ ,  $B=3000$ ,  $\text{initial}=5000$   
 $\text{Rs}=500/-$  from  $A \rightarrow B$   
 $T_1$   
 $R(A, a)$   
 $a=500$   
 $W(A, a) < -$  **Power failure**  $A=1500$ ,  $B=3000$ ,  $\text{final}=4500$   
 $\text{Read}(B, b)$   
 $b=b+500$   
 $W(B, b)$



# ACID Properties of Transaction

Transaction  
Manage-  
ment/ACID in  
DBMS

Dr. Munesh  
Singh

## Consistency

- **Consistency: (correctness)** Ensure that if the database was in a consistent state before the start of a transaction, then or termination, the database will also be in a consistent state.
- Example:  $\text{Sum}(A,B)$  before transaction =  $\text{Sum}(A,B)$  after transaction



# ACID Properties of Transaction

## Isolation

- **Isolation** Indicate that the actions performed by a transaction will be isolated or hidden from outside the transaction until it terminates
- Example: Given  $A=2000$ ,  $B=3000$ ,  $\text{initial}=5000$   
 $\text{Rs}=500/-$  from  $A \rightarrow B$

T1	T2
R(A,a)	
a-500	
W(A,a)	
	<b>4500</b>
	R(A,a)
	R(B,b)
	print(sum (a,b))
R(B,b)	
b=b+500	
W(B,b)	



# ACID Properties of Transaction

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

## Durability

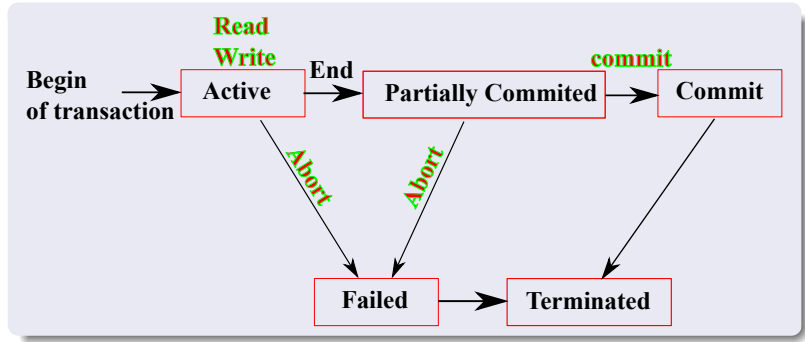
- **Durability** All updates done by a transaction must become permanent
- Ensure that the commit action of a transaction on its termination will be reflected in the database



# State of a Transaction

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh







# Concurrent Execution

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

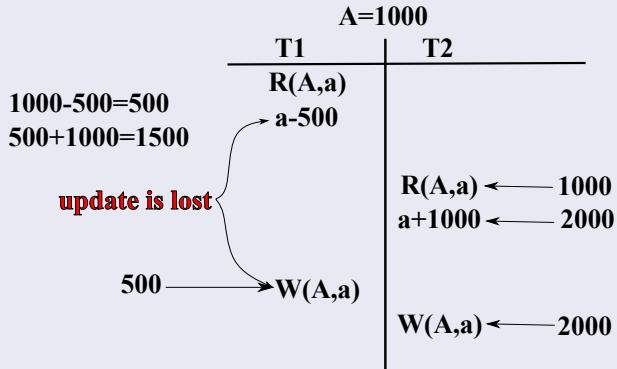
- It implies interleaving execution of operations of a transaction
- Benefits
  - Helps in reducing waiting time
  - Improve throughput and resource utilization
- **Schedule** It represents the order in which instructions of a transactions are executed



# Problem with Concurrent Execution

## Lost update Problem (W-W Conflict)

- Occurs when two transactions access the same database item have their operations interleaved in a way that makes the value of the database item incorrect.






# Problem with Concurrent Execution

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

## Temporary update (dirty read) Problem (W-R Conflict)

- Occurs when one transaction updates a database item and then the transaction fails, but its update is read by some other transaction.






# Problem with Concurrent Execution

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

## Unrepeatable Read (W-R Conflict)

- If a transaction 'T<sub>i</sub>' reads an item value twice and the item is changed by another transaction 'T<sub>j</sub>' in between the two read operation. Hence 'T<sub>i</sub>' receives different values for its two read operation of the same item






# Problem with Concurrent Execution

## Incorrect Summary Problem:

- If one transaction is calculating an aggregate summary function on a no of records, while other transaction is updating some of these records, the aggregate function may calculate some values before they are updated and other after they are updated results in incorrect summary.





# Schedules:

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

- A schedule 'S' of n transaction ' $T_1, T_2, \dots, T_n$ ' is an ordering of operations of the transactions in chronological order
- When several transactions are executing concurrently then the order of execution of various instructions is known as schedule



# Schedules:

## Types of schedule

- **Serial Schedule:** Does not interleave the actions of any operations of different transactions.
- Always ensure a consistent state

$T1 -> T2 -> T3$

$T1 -> T3 -> T2$

A=100	
T1	T2
R(A)	
A=A+50	
W(A)	
	R(A)
	A=A+100
	W(A)

Execution sequence for T1->T2:

100
150
150
250

Execution sequence for T2->T1:

100
200
200
250


A=100	
T1	T2
R(A)	
A=A+50	
	R(A)
	A=A+100
W(A)	
	W(A)



# Schedules:

## Types of schedule

- **Complete Schedule:** If the last operation of each transaction is either abort or commit.








## Types of schedule

- Recoverable Schedule:** Is one where for each pair of transactions  $(T_i, T_j)$ , such that  $T_j$  reads a data item that was previously written by  $T_i$ , then the commit operation of  $T_i$  should appear before commit operation of  $T_j$ .





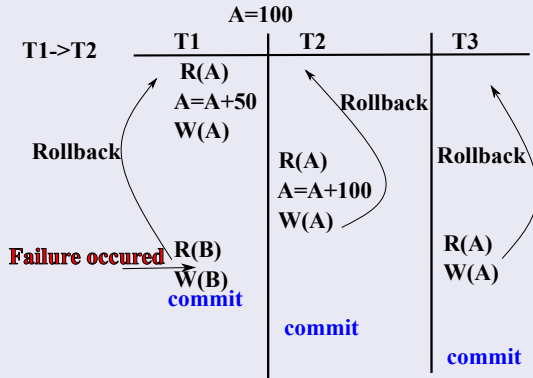
# Schedules:

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

## Types of schedule

- **Cascading rollback Schedule:**





# Schedules:

## Types of schedule

- **Cascadingless Schedule:** Is one where for each pair of transaction( $T_i, T_j$ ) such that  $T_j$  read a data item written by  $T_i$ , then the commit operation of  $T_i$  should appear the read operation of  $T_j$ .

T1	T2	T3
<div><div>Rollback</div><div><math>R(A)</math></div><div><math>A = A + 50</math></div><div><math>W(A)</math></div><div>commit</div></div>	<div><div><math>R(A)</math></div><div><math>A = A + 100</math></div><div><math>W(A)</math></div><div>commit</div></div>	<div><div><math>R(A)</math></div><div><math>W(A)</math></div><div>commit</div></div>



# Schedules:

## Types of schedule

- **Strict Schedule:** If a value written by a transaction cannot be read or overwritten by another transaction until the transaction is either aborted or committed
- Every strict schedule is both Recoverable and Cascade-less

T1	T2
R(A)	
A=A+50	
W(A)	
commit	
	R(A)
	A=A+100
	W(A)
	commit



# Conflict Transactions

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

## Conflict Operations

- Operations are said to be conflicting if:
  - Belong to different transactions.
  - Access to same database item 'A'
  - At-least one of them is a write operation. (R-W)(W-R)(W-W)
- **Equivalent Schedule** Two schedules 'S1' and 'S2' are said to be equivalent schedule if they produce the same final database state.
- **Result Equivalent Schedule:-** Produce same final database state for same initial value of data.



# Conflict Equivalent

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

- Two schedules are said to be conflict equivalent if all conflicting operations in both the schedule must be executed in the same order
- Question 1** Check for conflict Equivalent
  - S1: R1(A), R2(B), W1(A), W2(B)
  - S2: R2(B), R1(A), W2(B), W1(A)

S1		S2	
T1	T2	T1	T2
R(A)	R(B)	R(A)	R(B)
W(A)	W(B)	W(A)	W(B)

$S1 \sqsubseteq S2$



# Conflict Equivalent

- **Question 2** Check for conflict Equivalent
  - S1: R1(A),W1(A) R2(B),W2(B),R1(B)
  - S2: R1(A), W1(A), R1(B),R2(B),W2(B)



# Conflict Equivalent

- **Question 2** Check for conflict Equivalent

- S1: R1(A),W1(A) R2(B),W2(B),R1(B)
- S2: R1(A), W1(A), R1(B),R2(B),W2(B)

S1		S2	
T1	T2	T1	T2
R(A)		R(A)	
W(A)		W(A)	
	R(B)	R(B)	
	W(B)		R(B)
R(B)			W(B)
S1 \ S2			





# Serializability

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

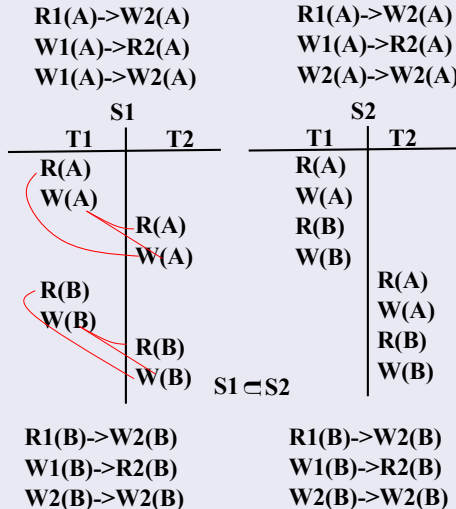
- A Schedule 'S' of 'n' transactions is serializable if it is equivalent to some serial schedule of the same 'n' transactions.
- **Conflict Serializable** If it is conflict equivalent to serial schedule



# Serializability Example

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh





# Serializability Gate Example

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

- S1: R1(X) R1(Y) R2(X) R2(Y) W2(Y) W1(X)
- S2: R1(X) R2(X) R2(Y) W2(Y) R1(Y) W1(X)
  - ① Both S1 and S2 are Conflict Serializable
  - ② Only S1 is Conflict Serializable
  - ③ Only S2 is conflict Serializable
  - ④ None



# Test for Conflict Serializability

## Precedence Graph is Used

- Let 'S' be a schedule, construct a directed graph known as precedence graph
- Graph consist of a pair of  $G=(V,E)$  where
  - V: is a set of vertices
  - E: set of edges
- **Algorithm for creation of graph**
  - Create a node for each transaction
  - A directed edge,  $T_i \rightarrow T_j$ , if  $T_i$  reads a value of an item written by  $T_i$ .
  - Directed edge  $T_i \rightarrow T_j$ , if  $T_j$  writes a value into item after it has been read by  $T_i$ .
  - Directed edge,  $T_i \rightarrow T_j$ , if  $T_j$  Write after  $T_i$  Write.
- A schedule is conflict serializable if and only if precedence graph is acyclic




# Test for Conflict Serializability

## Conflict Serializability

- **Question** Check for conflict Serializability

T1	T2	T3
R(X)		
R(Z)	R(Z)	
		R(X) R(Y) W(X)
	R(Y) W(Z) W(Y)	





# Test for Conflict Serializability


Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

## Conflict Serializability

- **Question** Check for conflict Serializability

T1	T2	T3
R(X)		
		R(X) W(X)
W(x)		
	R(X)	





# View Serializability

- Two schedules 'S' and 'S'' are view equivalent if the following conditions are met:
  - 1 For each data item Q, if  $T_i$  reads an initial value of in schedule S, then  $T_i$  in  $S'$  also reads an initial value of Q.
  - 2 If  $T_i$  executes Reads Q in S, and that value was produced by  $T_i$  (if any), then  $T_i$  must in schedule  $S'$  also reads the value of Q that was produced by  $T_j$ .
  - 3 For each data item Q, the transaction that perform the final write(Q) operation in schedule S must also perform the final write (Q) in schedule  $S'$ .
- A schedule is view serializable if it is view equivalent to a serial schedule.
- **Note:** Every conflict serializable schedule is also view serializable but not vice-versa,



# View Serializability Example

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

S1		S2	
T1	T2	T1	T2
R(A)		R(A)	
W(A)		W(A)	
	R(A)	R(B)	
	W(A)	W(B)	
			R(A)
R(B)			W(A)
W(B)			R(B)
	R(B)		W(B)
	W(B)		

$S1 \stackrel{v}{=} S2$





# Concurrency Control

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

- It is the process of managing simultaneous execution of transactions in a shared database, to ensure the serializability of transactions.
- **Purpose of concurrency control**
  - ① To enforce isolation
  - ② To preserve database consistency
  - ③ To resolve read-write and write-write conflicts



# Concurrency Control Techniques

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

- **Lock-Based Protocol:** A lock guarantees exclusive use of a data item to a current transaction.
  - 1 To Access Data item (Lock Acquired)
  - 2 After Completion of transaction (Release the Lock)
  - 3 All data items must be accessed in a mutually exclusive manner
- **Types of Locks:**
  - 1 Shared Lock (Lock-S): It is used for Read data item value
  - 2 Exclusive Lock (Lock-X) It is used for Both Read and Write



# Compatibility B/W Lock Model

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

	LockS	LockX
LockS	✓	X
LockX	X	X

T1	T2
Lock-X(B) Read(B) B+50 Write(B) Unlock-X(B)	Lock-S(B) Read(B) Unlock-S(B)

**Note: Any Number of Transaction can Hold Shared Lock, but Exclusive Lock only by One transaction at a time**



# Two-Phase Locking Protocol

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

- Requires both locks and unlocks being done in two phases.
- **Phases**
  - Growing (Expanding) Phase: New Locks on items can be acquired
  - Shrinking phase: Existing locks releases, but no new lock can be acquired
- As soon as final lock as been acquired it goes into **Lock Point**
- Once the Lock point is acquired, it can goes to shrinking phase



# Two-Phase Locking Protocol

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

- 2PL protocol enforces Serializability but may reduce concurrency due to the following reasons
  - ① Holding Lock UN-necessarily
  - ② Locking too early
  - ③ Penalty to other transaction



# Variations of 2PL Locking Protocol

- ① **Conservative (static) 2PL:** Acquire all lock before it starts. Release all locks after commit E.g: Avoid cascading rollback, Deadlock free
- ② **Strict 2PL:** Exclusive lock can't be released until commit. Helps in cascade-less schedule
- ③ **Rigorous 2PL** Shared/Exclusive can't be released until commit. Avoid cascading rollback
- ④ However, the **strict 2PL** and **Rigorous 2PL** may lead to Deadlock




# Deadlock

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

- A system is in a deadlock state if there exists a set of transactions such that every transaction in the set is waiting for another transaction in the set



(x,y)

**Lock(x)**

**x+y**

(x,y)

**Lock(y)**

**y-x**




# Deadlock Detection

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

- Simple way to detect a state of deadlock is to draw wait-for graph
- $G(V,E)$ ,  $V$  = Nodes describing transaction,  $E$  = Directed Edge
- $T_i \rightarrow T_j$  (directed edge)  $T_i$  is waiting for a resource data item held by  $T_j$

Transactions	Data items	Lock Mode
T1	Q	Shared
T2	P	Exclusive
T3	Q	Shared
T4	P	Exclusive



If cycle not occur, it means deadlock free





# Techniques to Control Deadlock

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

## Deadlock Prevention

- This protocol ensure that the system will never enter into a deadlock state
  - 1 Mutual Exclusion
  - 2 Hold and Wait
  - 3 No preemption
  - 4 Circular wait



# Other Techniques in Deadlock Prevention

## Use of timestamps

- Assuming that  $T_i$  request a data item currently held by  $T_j$ 
  - Wait-Die scheme**

if  $T_s(T_i) < T_s(T_j)$  [ $T_i$  is older than  $T_j$ ]  
 $T_i$  is allowed to wait otherwise if  $T_i$  is younger than  $T_j$   
then abort  $T_i$  ( $T_i$  dies) and restart it later with same timestamps.
  - Wound-wait scheme**

if  $T_s(T_i) < T_s(T_j)$  [ $T_i$  is older than  $T_j$ ]  
Then abort  $T_j$  ( $T_i$  wounds  $T_j$ ) and restart it with same timestamps. Otherwise  $T_s(T_i) \geq T_s(T_j)$  —  $T_i$  is allowed to wait



# Other Techniques in Deadlock Prevention

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

## Use of timestamps

- Assuming that  $T_i$  request a data item currently held by  $T_j$

### ① Time-Out Based Scheme

Based on Lock-timeouts, A transaction that has requested a Lock waits for at most a specified amount of time. If the lock is not granted within that time, transaction is said to timeout and it rolls itself back and restarts.



# Starvation

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

- A transaction is starved if it cannot process for an indefinite period of time while other transactions in the system continue normally
- Other reason of starvation: Algorithm dealing with deadlock prevention select some transaction as victim repeatedly, thus causing it to abort and never finish execution
- To prevent starvation:
  - 1 Modify timestamps
  - 2 Increase the priority of transaction



# Deadlock Recovery

- Selection of victim  $T1 < - > T2$  (abort  $T1$ ) selection of victim based on minimum cost
  - 1 Length of transaction (younger)
  - 2 Data item used by transaction [less no of data items]
  - 3 Data items that are to be locked [More data items to lock]
  - 4 How many transaction to be rollback [min of rollback]
- Rollback
  - 1 Full rollback [starting points]
  - 2 Partial rollback [saved points, locked point]
- Starvation: [Take care while selecting a victim so that a victim does not get starved]



# Deadlock Recovery

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

## Shadow Paging

- Requires fewer disk access than do-log methods.
- Maintain two page Tables during the life cycle of transaction
- When transaction start both page tables are identical
- Shadow page table is never changed over duration of transaction
- Current page table may change during Write operation
- All input and output operations use the current page table to lock database on disk
- Store shadow page table in Nonvolatile storage
- **When transaction commit system write current page table to nonvolatile storage. The current page table then becomes new shadowed P.T**



# Deadlock Recovery

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

## Shadow Paging

- Advantages
  - 1 Log-records overhead is removed
  - 2 Faster Recovery
- Drawback of shadow paging
  - 1 Commit overhead
  - 2 Data fragmentation
  - 3 Garbage collection



# Log Based Recovery

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

- Log is the most commonly used structure for recording the database modification
- Update log has following fields:-
  - 1 Transaction identifier
  - 2 Data item identifier
  - 3 Old value (prior to write)
  - 4 New value (after write)
- Example of Log record
  - 1 { *T1 start* }
  - 2 { *T1, x2, 10, 15* }
  - 3 { *T1 commit* }
  - 4 { *T1 Abort* }





# Log Based Recovery

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

## Write Ahead Log Strategy

- Log is written before any update is made to the database
- Transaction is not allowed to modify the physical data base until the undo portion of log is written to stable storage
- **Operations in Recovery Procedure**
  - ① UNDO ( $T_i$ ) – > Restore the old value
  - ② REDO( $T_i$ ) – > Updates by new values



# Approaches in Transaction Recovery Procedure:

Transaction  
Management/ACID in  
DBMS

Dr. Munesh  
Singh

## Deferred Database Modification

- Transaction do not immediately update the physical database
- Only transaction log is updated
- Database is physically updates only after the transaction reaches its commit point

## Immediate Database Modification

- Database is immediately updated by the transaction operations during the execution of transaction even before it reaches commit point
- In case of transaction abort before it reaches to commit point, a rollback or UNDO operation need to be done to restore the database to its earlier consistent state