

CUDA: 1) Matrix Addition 2) Matrix multiplication column major order 3) Matrix Multiplication Block based approach. Use input as a larger double number (64-bit). Run experiment for Threads = {1,2,4,8,16,32,64,128,256,500} Estimate the parallelization fraction

```
!pip install git+git://github.com/andreinechaev/nvcc4jupyter.git
```

```
Collecting git+git://github.com/andreinechaev/nvcc4jupyter.git
  Cloning git://github.com/andreinechaev/nvcc4jupyter.git to /tmp/pip-req-build-
  Running command git clone -q git://github.com/andreinechaev/nvcc4jupyter.git /
Building wheels for collected packages: NVCCPlugin
  Building wheel for NVCCPlugin (setup.py) ... done
  Created wheel for NVCCPlugin: filename=NVCCPlugin-0.0.2-cp36-none-any.whl size
  Stored in directory: /tmp/pip-ephem-wheel-cache-qsuqf__f/wheels/10/c2/05/ca241
Successfully built NVCCPlugin
Installing collected packages: NVCCPlugin
Successfully installed NVCCPlugin-0.0.2
```

```
%load_ext nvcc_plugin
```

```
created output directory at /content/src
Out bin /content/result.out
```

```
%%cu
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <bits/stdc++.h>
#include <chrono>
using namespace std::chrono;
using namespace std;
__global__ void Mat_add(double A[], double B[], double C[],int N,int th) {
    int id = threadIdx.x;
    for(int i=id ; i<N*N ; i+=th){
        C[i] = A[i] + B[i];
    }
}

int main(int argc, char* argv[]) {
    double *A, *B, *C;
    double *d_A, *d_B, *d_C;
    int N=10000;
    size_t size = N*N*sizeof(double);

    A = (double *) malloc(size);
    B = (double *) malloc(size);
    C = (double *) malloc(size);

    cudaMalloc(&d_A, size);
```

```

    cudaMalloc(&d_B, size);
    cudaMalloc(&d_C, size);
    int i,j;
    for (i=0; i<N; ++i){
        for(j=0 ; j<N ; ++j){
            A[i*N+j] = rand()%100000 + (1.0/(rand()%1000));
            B[i*N+j] = rand()%100000 + (1.0/(rand()%1000));
        }
    }

    /* Copy matrices from host memory to device memory */
    cudaMemcpy(d_A, A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, B, size, cudaMemcpyHostToDevice);

    int tt[10] = {1,2,4,8,16,32,64,128,256,500};
    for(int t=0 ; t<10 ; ++t){
        auto start = high_resolution_clock::now();
        Mat_add<<<1, tt[t]>>>(d_A, d_B, d_C,N, tt[t]);
        auto stop = high_resolution_clock::now();
        auto duration = duration_cast<microseconds>(stop - start);
        // cout << "Time taken by function: "<< duration.count() << " microseconds" << endl;
        cout <<duration.count()<<endl;
    }

    /* Wait for the kernel to complete */
    cudaThreadSynchronize();
    cudaMemcpy(C, d_C, size, cudaMemcpyDeviceToHost);

    /*for (int i=0; i<N; ++i){
        for(int j=0 ; j<N ; ++j){
            cout<<A[i*N + j]<<" + "<<B[i*N + j]<<" = "<<C[i*N + j]<<endl;
        }
    }*/
    cudaFree(d_A); cudaFree(d_B); cudaFree(d_C);
    free(A); free(B); free(C);

    return 0;
}

```

36

9

5

5

5

4

4

5

5

5

%%cu

```

#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include<bits/stdc++.h>
#include<chrono>
#define N 1000
using namespace std::chrono;
using namespace std;
__global__ void Mat_mul(double A[], double B[], double C[],int th) {
    int idx = threadIdx.x;

    for(int i = idx ; i<N*N ; i+=th){
        int temp=0;
        int y = i%N;
        int x = i/N;
        for(int k=0 ; k<N ; k++){
            temp += A[x*N+k] * B[k*N+y];
        }
        C[x*N+y] = temp;
    }
}

int main(int argc, char* argv[]) {
    double *A, *B, *C;
    double *d_A, *d_B, *d_C;
    //int N=4;
    size_t size = N*N*sizeof(double);

    A = (double *) malloc(size);
    B = (double *) malloc(size);
    C = (double *) malloc(size);

    cudaMalloc(&d_A, size);
    cudaMalloc(&d_B, size);
    cudaMalloc(&d_C, size);

    for (int i=0; i<N; ++i){
        for(int j=0 ; j<N ; ++j){
            A[i*N+j] = rand()%100000 + (1.0/(rand()%1000));
            B[i*N+j] = rand()%100000 + (1.0/(rand()%1000));
            /*A[i*N+j] = rand()%10;
            B[i*N+j] = rand()%10;*/
        }
    }
    /*cout<<"matrix A\n";
    for(int i=0 ; i<N ; ++i){
        for(int j=0 ; j<N ; ++j){
            cout<<A[i*N+j]<<" ";
        }
        cout<<endl;
    }
    cout<<endl;
    */
}

```

```

    cout<<endl;
    cout<<"matrix B\n";
    for(int i=0 ; i<N ; ++i){
        for(int j=0 ; j<N ; ++j){
            cout<<B[i*N+j]<<" ";
        }
        cout<<endl;
    }
    cout<<endl;*/

    /* Copy matrices from host memory to device memory */
    cudaMemcpy(d_A, A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, B, size, cudaMemcpyHostToDevice);

    int tt[10] = {1,2,4,8,16,32,64,128,256,500};
    for(int t=0 ; t<10 ; ++t){
        auto start = high_resolution_clock::now();
        Mat_mul<<<1, 2>>>(d_A, d_B, d_C, 2);
        auto stop = high_resolution_clock::now();
        auto duration = duration_cast<microseconds>(stop - start);
        cout << "Time taken by function: "<< duration.count() << " microseconds" << endl;
        //cout <<duration.count()<<endl;
    }

    /* Wait for the kernel to complete */
    //cudaThreadSynchronize();
    cudaMemcpy(C, d_C, size, cudaMemcpyDeviceToHost);

    /*cout<<"matrix C\n";
    for(int i=0 ; i<N ; ++i){
        for(int j=0 ; j<N ; ++j){
            cout<<C[i*N+j]<<" ";
        }
        cout<<endl;
    }*/
    cudaFree(d_A); cudaFree(d_B); cudaFree(d_C);
    free(A); free(B); free(C);

    return 0;
}

```

```

↳ Time taken by function: 27 microseconds
Time taken by function: 9 microseconds
Time taken by function: 4 microseconds
Time taken by function: 5 microseconds
Time taken by function: 5 microseconds
Time taken by function: 5 microseconds
Time taken by function: 4 microseconds
Time taken by function: 4 microseconds
Time taken by function: 4 microseconds
Time taken by function: 5 microseconds

```

